

# Suitability analysis of Kubernetes for Seznam.cz

Bc. Ondřej Šejvl

Department of Software Engineering  
Faculty of Information Technology  
Czech Technical University in Prague

June 13, 2016



# Content

- ① About
- ② Tarsier
- ③ Conclusion
- ④ Discussion

# Assignment

- Make yourself familiar with the Kubernetes system.
- Compare Kubernetes in different environments.
- Design and implement a set of applications to simulate real traffic.
  - Load configuration files, SSL certificates and static resources.
  - Use persistent storage.
  - Raise controlled faults, deadlocks, segmentation faults.
  - Simulate heavy CPU and RAM load.
  - Log their own activity and reliably transfer the logs to the central log storage.

# Technologies?

## Kubernetes

is an open-source system for automating deployment, operations, and scaling of containerized applications. [1]



Figure: Kubernetes [2]

## Docker

allows you to package an application with all of its dependencies into a standardized unit for software development. [3]



Figure: Docker [4]

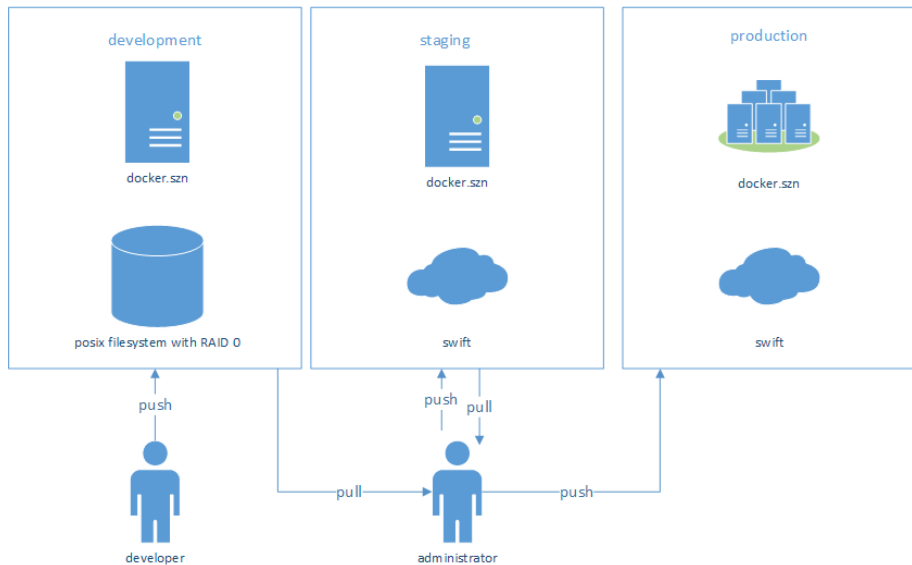
# Possible problems with Kubernetes in Seznam.cz

- Docker registry
- Secrets distribution
- Logging
- Security
- Monitoring
- Static content of websites



Figure: Krasty [5]

# Docker registry



# Logging I

## Debug logs

- textual
- no need for 100% reliability

## Business logs

- structured – Apache Avro [6]
- 100% reliability
- not part of the thesis

## Third-party logs

- system logs
- third-party applications (nginx mostly, ...)

**Goal:** One solution to rule them all.

# Logging II

## Kafkalog

- library for storing, rotating, compressing logs

## Kafkafeeder

- wrapper around Heka [7]
- sending messages from Kafkalog to Kafka [8]
- with acknowledgement



# Logging III

## Debug logs

- Kafkalog handler to native loggers (Go Logrus [9])

## Business logs

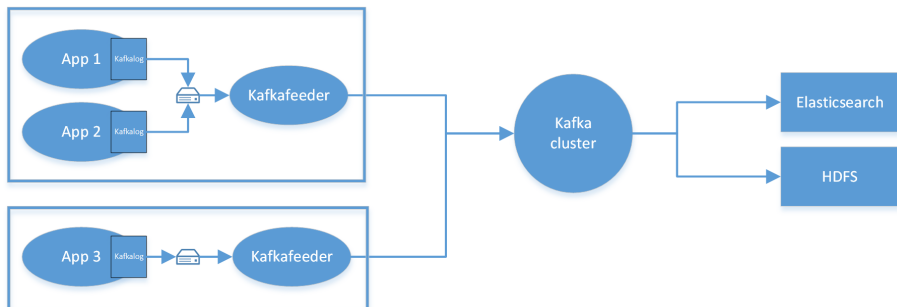
- additional library that uses Kafkalog
- not part of the thesis

## Third-party logs

- simple application with various input interfaces (stdin, syslog, ...)
- save to Kafkalog
- not part of the thesis

# Logging IV

Machine1



Machine 2

Figure: Logging flow

# Tarsier I

- Testing application
- Modular
  - Heavy load – Gobble RAM, Spin CPU
  - Faults – Deadlock, Segfault
  - Net – Dial
  - Persistent storage – Open, Read, Write
  - Sleeping beauty – Sleep



Figure: Tarsier [10]

# Tarsier II

```
curl -i -X POST -d '
    command: "persistent_storage/write"
    data:
        amount: "100_MB"
        files: 10
    wave:
        remains: 3
        buddies: 3
' kubernetes.dev:6666/exec
```

## Tarsier III




Figure: Tarsier wave

# Output

<https://github.com/sejvlond/master-thesis>

- Tarsier** Application for testing Kubernetes cluster
- Kafkafeeder** Application that transfers logs to Kafka
- Go kafkalog** Apache Kafka log format implementation in Go
- Kafkalog Logrus** Kafkalog hook to logrus
  - Heka** Some Heka adjustment for greater reliability
- Heka Kafkalog** Heka Splitter and Decored plugins that use go-kafkalog
- Kubernetes testing** Support scripts for testing Kubernetes
- Go Ultimate server** Simple Go server for benchmarking
- Python Ultimate server** Simple Python Tornado server for benchmarking

# Benefits

- ① Kubernetes
- ② Docker
- ③ Golang
- ④ Kafka
- ⑤ Heka
- ⑥ Fluentd
- ⑦ Logstash

# Acknowledgements






- Ing. Jan Baier
- Ing. Tomáš Kukrál
- David Bouček
- Martin Stružský
- Seznam.cz, a.s.








Figure: Krasty [5]



# References I

-  Google, Inc. *Kubernetes*. [online] [cit. 2016-06-10]. Available from: <http://kubernetes.io/>
-  Twitter, Inc. *Kubernets logo*. [online] [cit. 2016-06-10]. Available from: [https://pbs.twimg.com/profile\\_images/511909265720614913/21\\_d3cvM.png](https://pbs.twimg.com/profile_images/511909265720614913/21_d3cvM.png)
-  Docker, Inc. *What is Docker?* [online] [cit. 2016-06-10]. Available from: <https://www.docker.com/what-docker>
-  Xebialabs, Inc. *Docker logo*. [online] [cit. 2016-06-10]. Available from: <http://blog.xebialabs.com/wp-content/uploads/2016/04/docker-logo.png>
-  Seznam.cz, a.s. *Krasty*. [online] [cit. 2016-06-10]. Available from: [krastyhttp://onas.seznam.cz/cz/ke-stazeni/krasty/](http://onas.seznam.cz/cz/ke-stazeni/krasty/)

# References II

-  The Apache Software Foundation. *Apache Avro*. [online] [cit. 2016-06-10]. Available from: <https://avro.apache.org/>
-  Mozilla Corporation. *Heka*. [online] [cit. 2016-06-10]. Available from: <https://github.com/mozilla-services/heka>
-  The Apache Software Foundation. *Apache Kafka*. [online] [cit. 2016-06-10]. Available from: <http://kafka.apache.org/>
-  Eskildsen S. *Logrus*. [online] [cit. 2016-06-10]. Available from: <https://github.com/Sirupsen/logrus>
-  The Animal Network. *Tarsier*. [online] [cit. 2016-06-10]. Available from: <http://animalnetwork.davyprovoost.netdna-cdn.com/wp-content/uploads/2015/02/tarsier-funny-600x380.jpg>

# References III



Docker, Inc. *Configure logging drivers*. [online] [cit. 2016-06-10].

Available from:

<https://docs.docker.com/engine/admin/logging/overview/>



Treasure Data, Inc. *Fluentd architecture*. [online] [cit. 2016-06-10].

Available from:

<http://docs.fluentd.org/images/fluentd-architecture.png>

**Thank you for your attention**

**Bc. Ondřej Šejvl**  
sejvlond@fit.cvut.cz

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file.

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog.



## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald.

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.

~~Graylog cluster, Logstash~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.

~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input).

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.

~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs.



## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

**splunk** to splunk using HTTP Event Collector.

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

**splunk** to splunk using HTTP Event Collector. ~~Splunk Cloud~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

**splunk** to splunk using HTTP Event Collector. ~~Splunk Cloud~~

**etwlogs** ETW logging driver on Windows.

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

**splunk** to splunk using HTTP Event Collector. ~~Splunk Cloud~~

**etwlogs** ETW logging driver on Windows. ~~Windows~~

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

**splunk** to splunk using HTTP Event Collector. ~~Splunk Cloud~~

**etwlogs** ETW logging driver on Windows. ~~Windows~~

**gcplogs** to Google Cloud Logging.

## Question

Why did you chose Kafkafeeder instead of native Docker's log drivers?

## Answer

Docker's log drivers for container standard output [11]

**json-file** JSON messages to file. ~~Kafka~~

**syslog** to syslog. ~~Kafka, 1 kB~~

**journald** to journald. ~~systemd~~

**gelf** to a GELF endpoint like Graylog or Logstash.  
~~Graylog cluster, Logstash~~

**fluentd** to fluentd (forward input). **possible for debug logs**

**awslogs** to Amazon CloudWatch Logs. ~~Amazon CloudWatch~~

**splunk** to splunk using HTTP Event Collector. ~~Splunk Cloud~~

**etwlogs** ETW logging driver on Windows. ~~Windows~~

**gcplogs** to Google Cloud Logging. ~~Google Cloud~~

## fluentd

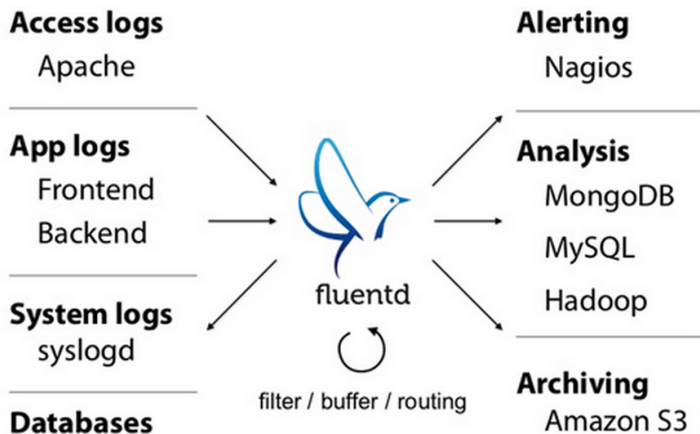


Figure: fluentd [12]



**Thank you for your attention**

**Bc. Ondřej Šejvl**  
sejvlond@fit.cvut.cz